

# Database Management System 22

## Two-Phase Locking & Timestamp Based Protocols

### Two-Phase Locking Protocol

Variations of Two-Phase  
Locking

Lock Conversions

### Timestamp-Based Protocols

Timestamps

Timestamp Ordering  
Protocols

Problems with Timestamp  
Ordering Protocols

Thomas' Write Rule

Chittaranjan Pradhan  
School of Computer Engineering,  
KIIT University

## Two-Phase Locking Protocol

Two-phase locking protocol is a protocol which ensures serializability

This protocol requires that each transaction issues lock and unlock requests in two phases. The two phases are:

- **Growing phase:** Here, a transaction acquires all required locks without unlocking any data, i.e. the transaction may not release any lock
- **Shrinking phase:** Here, a transaction releases all locks and cannot obtain any new lock

The point in the schedule where the transaction has obtained its final lock is called the **lock point** of the transaction

*Transactions can be ordered according to their lock points*

### Two-Phase Locking Protocol

Variations of Two-Phase Locking

Lock Conversions

### Timestamp-Based Protocols

Timestamps

Timestamp Ordering Protocols

Problems with Timestamp Ordering Protocols

Thomas' Write Rule

# Two-Phase Locking Protocol...

## Two-Phase Locking Protocol

Variations of Two-Phase Locking

Lock Conversions

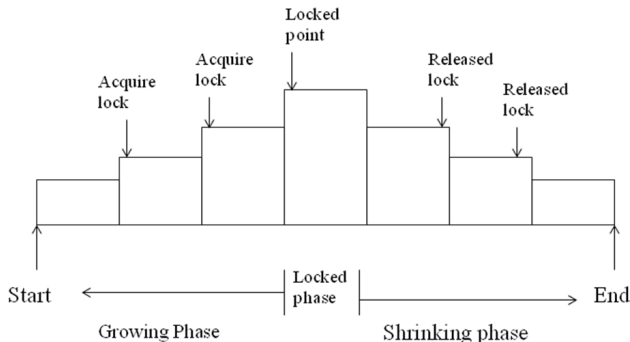
## Timestamp-Based Protocols

Timestamps

Timestamp Ordering Protocols

Problems with Timestamp Ordering Protocols

Thomas' Write Rule



- Two transactions cannot have conflicting locks
- No unlock operation can precede a lock operation in the same transaction
- No data are affected until all locks are obtained, i.e. until the transaction is in its locked point
- Two-phase locking may limit the amount of concurrency that can occur in a schedule

## Variations of Two-Phase Locking

### Conservative or Static 2PL

- Lock all the data items that you need before transaction starts execution by predeclaring its *read-set* and *write-set*
- Conservative 2PL is a deadlock-free protocol
- It is difficult to implement because of the need to predeclare the read-set and write-set

### Strict 2PL

- A transaction T doesn't release any of its exclusive locks until after it commits or aborts
- No other transaction can read/write an item that is written by T unless T has committed
- Strict 2PL is not deadlock-free

# Variations of Two-Phase Locking...

## Variations of Two-Phase Locking...

### Rigorous 2PL

- A transaction T doesn't release any of its locks until after it commits or aborts
- Strict 2PL holds write-locks until it commits; whereas Rigorous 2PL holds all locks
- Conservative 2PL must lock all its data items before it starts, so once the transaction starts it is in shrinking phase; whereas Rigorous 2PL doesn't unlock any of its data items until after it terminates

# Lock Conversions

| $T_7$      |
|------------|
| Read(a1);  |
| Read(a2);  |
| Read(a3);  |
| ...        |
| Read(an);  |
| Write(a1); |

| $T_8$           |
|-----------------|
| Read(a1);       |
| Read(a3);       |
| Display(a1+a3); |

## Lock Conversions

**Lock Upgrade:** This is the process in which a shared lock is upgraded to an exclusive lock

**Lock Downgrade:** This is the process in which an exclusive lock is downgraded to a shared lock

*Lock upgrading can take place only in the growing phase, where as lock downgrading can take place only in the shrinking phase*

Thus, the two-phase locking protocol with lock conversions:

- **First Phase:**
  - Can acquire a lock-S on item
  - Can acquire a lock-X on item
  - Can convert a lock-S to a lock-X (upgrade)
- **Second Phase:**
  - Can release a lock-S
  - Can release a lock-X
  - Can convert a lock-X to a lock-S (downgrade)

Like the basic two-phase locking protocol, two-phase locking with lock conversion generates only conflict-serializable schedules and transactions can be serialized by their lock points

If the exclusive locks are held until the end of the transaction, then the schedules became cascadeless

## Lock Conversions...

Schedule8

| $T_7$   | $T_8$   |
|---|---|
| Lock-S(a1);<br>Read(a1);<br><br>Lock-S(a2);<br>Read(a2);<br>Lock-S(a3);<br>Read(a3);<br><br>...<br>Lock-S(an);<br>Read(an);<br>Upgrade(a1);<br>Write(a1); | Lock-S(a1);<br>Read(a1);<br><br><br><br><br><br>Lock-S(a3);<br>Read(a3);<br>Display(a1+a3);<br>Unlock(a1);<br>Unlock(a3); |



# Timestamp-Based Protocols

## Timestamp-Based Protocols

The timestamp method for concurrency control doesn't need any locks and therefore this method is free from deadlock situation

Locking methods generally prevent conflicts by making transaction to wait; whereas timestamp methods do not make the transactions to wait. Rather, transactions involved in a conflicting situation are simply rolled back and restarted

A **timestamp** is a unique identifier created by the Database system that indicates the relative starting time of a transaction. Timestamps are generated either using the system clocks or by incrementing a logical counter every time a new transaction starts

Timestamp protocol is a concurrency control protocol in which the fundamental goal is to order the transactions globally in such a way that older transactions get priority in the event of a conflict

## Timestamps

Timestamp  $TS(T_i)$  is assigned by the database system before the transaction  $T_i$  starts its execution

The timestamps of the transactions determine the serializability order. Thus, if  $TS(T_i) < TS(T_j)$ , then the system must ensure that the produced schedule is equivalent to a serial schedule in which  $T_i$  appears before  $T_j$

There are two timestamp values associated with each data item Q:

- **W-Timestamp(Q):** It denotes the largest timestamp of any transaction that executed write(Q) operation successfully
- **R-Timestamp(Q):** It denotes the largest timestamp of any transaction that executed read(Q) operation successfully

Two-Phase Locking  
Protocol

Variations of Two-Phase  
Locking

Lock Conversions

Timestamp-Based  
Protocols

Timestamps

Timestamp Ordering  
Protocols

Problems with Timestamp  
Ordering Protocols

Thomas' Write Rule

## Timestamp Ordering Protocols

This ensures that any conflicting read and write operations are executed in timestamp order

- Suppose transaction  $T_i$  issues read(Q):
  - If  $TS(T_i) < W-TS(Q)$ , then  $T_i$  needs to read a value of Q that was already overwritten. Hence, the read operation is rejected and  $T_i$  is rolled back
  - If  $TS(T_i) \geq W-TS(Q)$ , then the read operation is executed, and  $R-TS(Q)$  is set to maximum of  $R-TS(Q)$  and  $TS(T_i)$
- Suppose transaction  $T_i$  issues write(Q):
  - If  $TS(T_i) < R-TS(Q)$ , then the value of Q that  $T_i$  is producing was needed previously, and the system assumed that the value would never be produced. Hence, the system rejects the write operation and rolls  $T_i$  back
  - If  $TS(T_i) < W-TS(Q)$ , then  $T_i$  is attempting to write an obsolete value of Q. Hence, the system rejects this write operation and rolls  $T_i$  back
  - Otherwise, the system executes the write operation and sets  $W-TS(Q)$  to  $TS(T_i)$

# Timestamp Ordering Protocols...

If a transaction  $T_i$  is rolled back by the concurrency-control scheme, the system assigns it a new timestamp and restarts it

| $T_1$  | $T_2$  | $T_1$  | $T_2$  | $T_1$  | $T_2$  |
|--|--|--|--|--|--|
| Read(A);<br>A:=A-100;<br>Write(A);<br><br>Read(B);<br>B:=B+100;<br>Write(B); | Read(A);<br>Temp:=0.2*A;<br>A:=A-Temp;<br>Write(A);<br><br>Read(B);<br>B:=B+Temp;<br>Write(B); | Read(A);<br>A:=A-100;<br><br>Write(A);<br>Read(B);<br>B:=B+100;<br>Write(B); | Read(A);<br>Temp:=0.2*A;<br>A:=A-Temp;<br>Write(A);<br>Read(B);<br><br>B:=B+Temp;<br>Write(B); | Read(A);<br>A:=A-100;<br>Write(A);<br><br>Read(B);<br>B:=B+100;<br><br>Write(B); | Read(A);<br>Temp:=0.2*A;<br>A:=A-Temp;<br><br>Write(A);<br>Read(B);<br>B:=B+Temp;<br>Write(B); |
| Schedule3  |  | Schedule4  |  | Schedule5  |  |

## Timestamp Ordering Protocols...

From the above examples, we can say that the timestamp ordering protocol always ensures conflict serializability. This is because conflicting operations are processed in timestamp order

The protocol ensures freedom from deadlock, since no transaction ever waits

However, there is a possibility of starvation of long transactions if a sequence of conflicting short transactions causes repeated restarting of the long transaction

If a transaction is found to be getting restarted repeatedly; conflicting transactions need to be temporarily blocked to enable the transaction to finish

## Problems with Timestamp Ordering Protocols

Suppose  $T_i$  aborts, but  $T_j$  has read a data item written by  $T_i$ . Then,  $T_j$  must abort; if  $T_j$  had been allowed to omit earlier, the schedule is not recoverable. Further, any transaction that has read a data item written by  $T_j$  must abort. This can lead to **cascading rollback**; i.e. chain of rollbacks

The solutions to this type of problem are:

- A transaction is structured in such a way that its writes are performed at the end of its processing. At the same time, all writes of a transaction form an atomic action; i.e. no transaction may execute while a transaction is being written, or
- Wait for the data to be committed before reading it

## Thomas' Write Rule

One of the problems with basic timestamp ordering protocol is that it does not guarantee recoverable schedules

A modification to the basic timestamp ordering protocol that relaxes the conflict serializability can be used to provide greater concurrency by rejecting obsolete write operations

Schedule9

| $T_9$  | $T_{10}$  |
|--|-----------|
| Read(Q);<br><br>Write(Q);<br>Read(M);<br>Write(M); | Write(Q); |

### Thomas' Write Rule...

- $T_i$  issues Read(Q):
  - Remains same
- $T_i$  issues write(Q):
  - If  $TS(T_i) < R-TS(Q)$ ,  $T_i$  is aborted, rolled back and is restarted with a new timestamp
  - If  $TS(T_i) < W-TS(Q)$ ,  $T_i$  is attempting to write an obsolete value of Q. Hence, this write operation can be ignored
  - Otherwise, the system executes the write operation and sets  $W-TS(Q) = TS(T_i)$

Thomas' write rule makes use of view Serializability by deleting obsolete write operations from the transactions that issue them